

K-Nearest Neighbours (KNN)

KNN is a supervised machine learning algorithm that can be used for both classification and regression tasks. It is a simple yet effective algorithm based on the principle that similar things are grouped together.

→ **How it works:**

1. **Training phase** → The algorithm stores all the training data points without any learning process.

2. **Prediction phase** → When a new data point (query point) needs to be classified or predicted →

(a) It calculates distance between all the training point and query point.

(b) Find the k -nearest neighbours to the query point.

(c) If the task is classification, assigns the query point to the most frequent class among the k neighbours.

(d) If the task is regression, calculate the average of the target values of the k neighbours and assign it to the query point.

* **Distance Metric** : The closeness between data points is measured using a distance metric. The most common metric used is **euclidean distance** which is given by the formula shown below and calculates straight line distance.

$$d(x, y) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

Suppose two points in space are (3, 4) and (6, 3) then distance between them using euclidean distance is—

$$d = \sqrt{(3-6)^2 + (4-3)^2} = \sqrt{9+1} = \sqrt{10} \approx 3.16$$

Manhattan distance → The manhattan distance also known as (L1 distance or taxicab distance) between two points is the sum of the absolute differences of their coordinates. It is calculated using the formula—

$$d_{\text{manhattan}}(p, q) = \sum_{i=1}^n |p_i - q_i|$$

where p & q are two points in an

n -dimensional space, and p_i and q_i are the coordinates of these points in the i^{th} dimension.

Manhattan Distance

The Manhattan distance measures how far two points are by moving along the gridlines (like a taxi navigating the streets of a city). The distance metric is more appropriate in situations where we can only move along the gridlines rather than in a straight line.

The Manhattan distance between two points (3,4) & (6,3) can be calculated as—

$$d = |6-3| + |3-4| = 3 + 1 = 4$$

- It can be used in places where we require grid-based movement (e.g. city blocks, pixel grids in images). In case of high dimensional spaces and sparse data.

Minkowski Distance: Minkowski is a generalization of both Euclidean and Manhattan distances. It is calculated using the formula:

$$d_{\text{Minkowski}}(p, q) = \left(\sum_{i=1}^n |p_i - q_i|^p \right)^{1/p}$$

where p is the parameter that defines the type of distance. When $p=1$, Minkowski distance becomes Manhattan distance and when $p=2$, Minkowski distance becomes Euclidean distance. When p increases, the distance metric becomes more sensitive to large differences better in any single dimension.

- If we suspect that some features might disproportionately influence the distance, we can choose a higher p to emphasize these differences.

The Minkowski distance between two points (3,4) & (6,3) taking $p=3$ & $p=4$, can be calculated as—

For $p=3$, $d = (|6-3|^3 + |3-4|^3)^{1/3} = (3^3 + 1^3)^{1/3} = 28^{1/3}$
 $d \approx 3.03$

For $p=4$, $d = (|6-3|^4 + |3-4|^4)^{1/4} = (3^4 + 1^4)^{1/4} = 82^{1/4}$
 $d \approx 3$

For $p=10$, $d = (|6-3|^{10} + |3-4|^{10})^{1/10} = (3^{10} + 1^{10})^{1/10} = (59050)^{1/10}$
 ≈ 3

* Choosing the correct value of k : The choice of k in KNN significantly impacts the model's performance. A small k might lead to overfitting, while a large k can result in underfitting (ignoring local patterns).

- (a) Using cross-validation can be a suitable value method for choosing k .
- (b) We can even do a hyperparameter search using different values of k to find the best k .
- (c) We can consider square root of the dataset size as a starting point and start by choosing an odd value.
- (d) For imbalanced datasets, a smaller k might be more appropriate to avoid the majority class dominating the predictions.
- (e) If the data is noisy and contains many outliers, a larger k can help smooth out the noise and make the model more robust.

Pros and Cons of KNN algorithm:

- Pros:
- (a) It is simple and easy to understand.
 - (b) No training phase is required (a lazy learning algorithm).
 - (c) It can handle high-dimensional data effectively.
- Cons:
- (a) Can be computationally expensive for large datasets (because it needs to calculate distance to every data point), especially during prediction.
 - (b) It is sensitive to the choice of both ' k ' and distance metric.
 - (c) It can be affected by presence of noisy data and outliers.

Applications:

- (a) Recommendation Systems: Suggesting similar items based on user preferences. in image
- (b) Image Recognition: For identifying and classifying objects.
- (c) Anomaly detection: Identifying unusual patterns that do not confirm to expected behaviour.